

Understanding Fabric Capacities

Benni De Jagere



Slides



b3

CEΘAL

REBTECH



DATAmasterminds



Measure Killer



Snabbfoting

TIMEXTENDER



Telia



Benni De Jagere

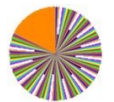
Senior Program Manager | Fabric Customer Advisory Team (FabricCAT)



dataMinds



sessionize



Fabric CAT

.be Member

@BenniDeJagere

/bennidejagere

/bennidejagere

/bennidejagere

#SayNoToPieCharts





Fabric Capacities Introduction



Microsoft Fabric



Data
Factory



Synapse Data
Engineering



Synapse Data
Science



Synapse Data
Warehouse



Synapse Real
Time Analytics



Power BI



Data
Activator

AI Assisted

Shared Workspaces

Universal Compute Capacities

One Security

OneLake

Intelligent data foundation

Single...

Onboarding and trials
Sign-on
Navigation model
UX model
Workspace organization
Collaboration experience
Data Lake
Storage format
Data copy for all engines
Security model
CI/CD
Monitoring hub
Governance & Capacity Metrics
Data Hub

Capacities are to Fabric what CPUs are to PCs

Personal Computing

When you purchase a PC you choose the number of CPU cores. The more CPU cores the more load the PC can handle.



The CPU cores are dynamically shared across all applications with no need to pre-allocate by app.



The total consumption of the CPU across all the apps cannot exceed the number of cores. CPU overload causes a slowdown.

Name	Status	48% CPU	82% Memory	1% Disk
Apps (10)				
> Task Manager		2.6%	65.0 MB	0 MB/s
> Snipping Tool (2)		0%	28.1 MB	0 MB/s
> Settings		0%	45.3 MB	0 MB/s
> Microsoft Word (2)		0.5%	131.3 MB	0 MB/s
> Microsoft Teams (9)		4.5%	800.3 MB	0.1 MB/s
> Microsoft PowerPoint (2)		1.5%	300.5 MB	0.1 MB/s
> Microsoft Outlook (13)		0.9%	414.0 MB	0.1 MB/s
> Microsoft OneNote		0%	85.2 MB	0 MB/s
> Microsoft Excel (2)		0.8%	93.5 MB	0 MB/s
> Microsoft Edge (19)		0%	1,628.3 MB	0.1 MB/s

Fabric Capacities

In Fabric, you provision a Capacity with a number of "capacity units". The more capacity units provisioned, the more compute load handled. Unlike the PC, capacity units can be scaled up or down as needed.

The capacity units are dynamically shared across all the Fabric workloads, with no pre-allocation necessary. A single capacity can simultaneously drive BI, DW, Spark, ML and every other compute engine in Fabric

The total consumption of the capacity across all the workloads cannot exceed the capacity units provisioned. Overloading the capacity will throttle it (slow down). Auto scale can dynamically increase the available compute units avoiding the slowdown.

Capacities are a shared resource

Shared across workloads

A single capacity is providing the compute power for all Fabric workloads.


There is no need to allocate compute for each workload separately.



Shared Across Projects

A single capacity typically supports dozens of separate projects simultaneously, each managed in its own workspace.

It is rare to have a capacity dedicated to a single project




Shared across users


For each project, many developers will share a workspace where collaborative development and consumption at scale is managed.

Each creator can provision any artifact and run any job without the need for any pre-approval or planning

Developers/Creators



Consumers




Capacities are flexible building blocks for growth

Capacities can be configured in endless ways to meet scale, usage and governance requirements while tuning to minimize TCO and performance goals


Scale Vertically

Increased capacity size provides more throughput




Scale Horizontally

Scale horizontally using the benefits of modular design for hardened isolation and governance



Regional Availability

Use different capacities for different regions to support GDPR / Data residency requirements



Provisioning and Deploying Capacities

Purchased in Azure

- **Purchased** either as a PAYG or RI resource
- **Provisioned with a certain amount of compute** units, analogous to CPU cores.
- The **more capacity units are provisioned, the more load** the capacity can support
 - Multiply SKU size by 30s to match platform evaluation in metrics app
- Capacities are **priced at a fixed hourly rate**, based on capacity units provisioned
- The RI commitment (1-year reserved instance) enjoys a **40% discount**

Universal Compute Capacities SKU Sizing

SKU	Capacity Units (CU)	CU's (per 30s)	Power BI SKU	Power BI V-cores
F2	2	60	-	0.25
F4	4	120	-	0.5
F8	8	240	A1	1
F16	16	480	A2	2
F32	32	960	A3	4
F64	64	1920	P1	8
F128	128	3840	P2	16
F256	256	7680	P3	32
F512	512	15360	P4	64
F1024	1024	30720	P5	128
F2048	2048	61440	-	256

Provisioning and Deploying Capacities

Deployed to Regions

- Each capacity **resides in a specific region of the buyers' choice** where both the data & compute reside
- **Workspaces are assigned to a capacity** that provides the compute and storage for all the workspace artifacts
- Multiple capacities can be purchased, deployed and managed by **different owners** residing in a single tenant allowing each business unit to pay for their own consumption

Tenant (Microsoft)

Capacity 1 (West Europe)

Workspace 1 (Data Science Team)

Item 1
Power BI Dataset

Item 2
AI Function
Evaluation

Workspace 2 (Research)

Item 3
Spark Notebook

Item 4
ML Dataflow

Capacity 2 (Central India)

Capacity 3 (West US)



Bursting and Smoothing

Smoothing intro and benefits

Load stabilization

Smoothing helps capacities self-stabilize by flattening large spikey loads into a smooth load profile, eliminating temporal spikes

Eliminates Scheduling contention

Large/scheduled Jobs usage (not execution) are smoothed over 24 hours, eliminating the need to decide the timing and order of job execution

Performance now, pay later. Demo to follow.

Bad actor protection

Interactive operations smoothed over several minutes, preventing a single user with a very demanding query from hogging the entire capacity



What is Bursting?

Job acceleration

Bursting provides extra compute resources to jobs and queries to accelerate their completion

Go beyond

The extra resources of bursting allow jobs to **utilize far more resources than "face value"**

Instead of running a job on 64 CU and completing in 60 seconds, bursting could use 256 CUs to complete the job in 15 seconds.

Same amount of work, just completed faster

No hassle, No overload

Bursting is automatic when the system reasons it can accelerate the job by applying extra resources. No settings are required.

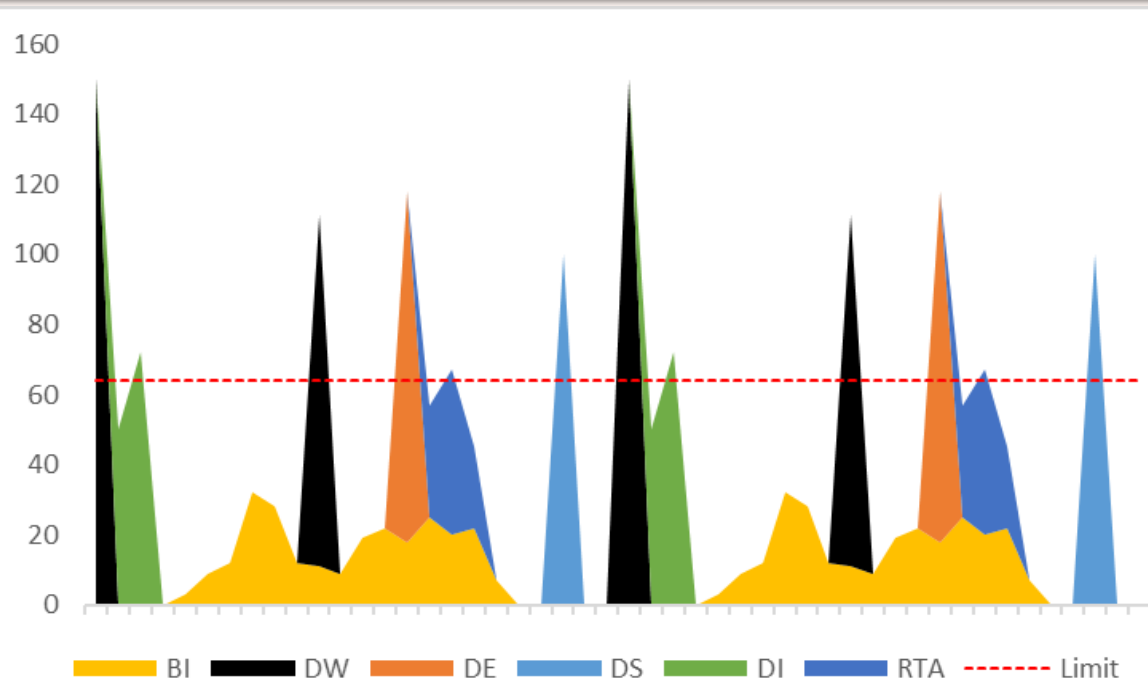
Bursting prevents an overload as the *smoothing* mechanism will always flatten the resource burst

Bursting and smoothing | before and after

Looking at an example of a 64 CU capacity, running multiple workloads over a couple of days...


Before Smoothing

- Actual load as executed on the capacity before smoothing
- *Bursting* accelerates jobs execution by resource boosting
- The capacity could be overloaded 25% of the time
- Some of the overloads are more than 2x the limit
- There are periods of no/low usage




After Smoothing

- Shows the reported load (not runtime execution) against the capacity limits
- There is NO overload, and consumption is more stable
- The smoothing of usage fills in gaps




Jobs Executed



Bursting and Smoothing


- Regardless of SKU, Fabric *bursting* will automatically allocate resources as needed to execute at maximum performance
- As such, one query could consume all the quota of a single time window and much more!
- To avoid an overload, *smoothing* kicks in

Bursting and Smoothing



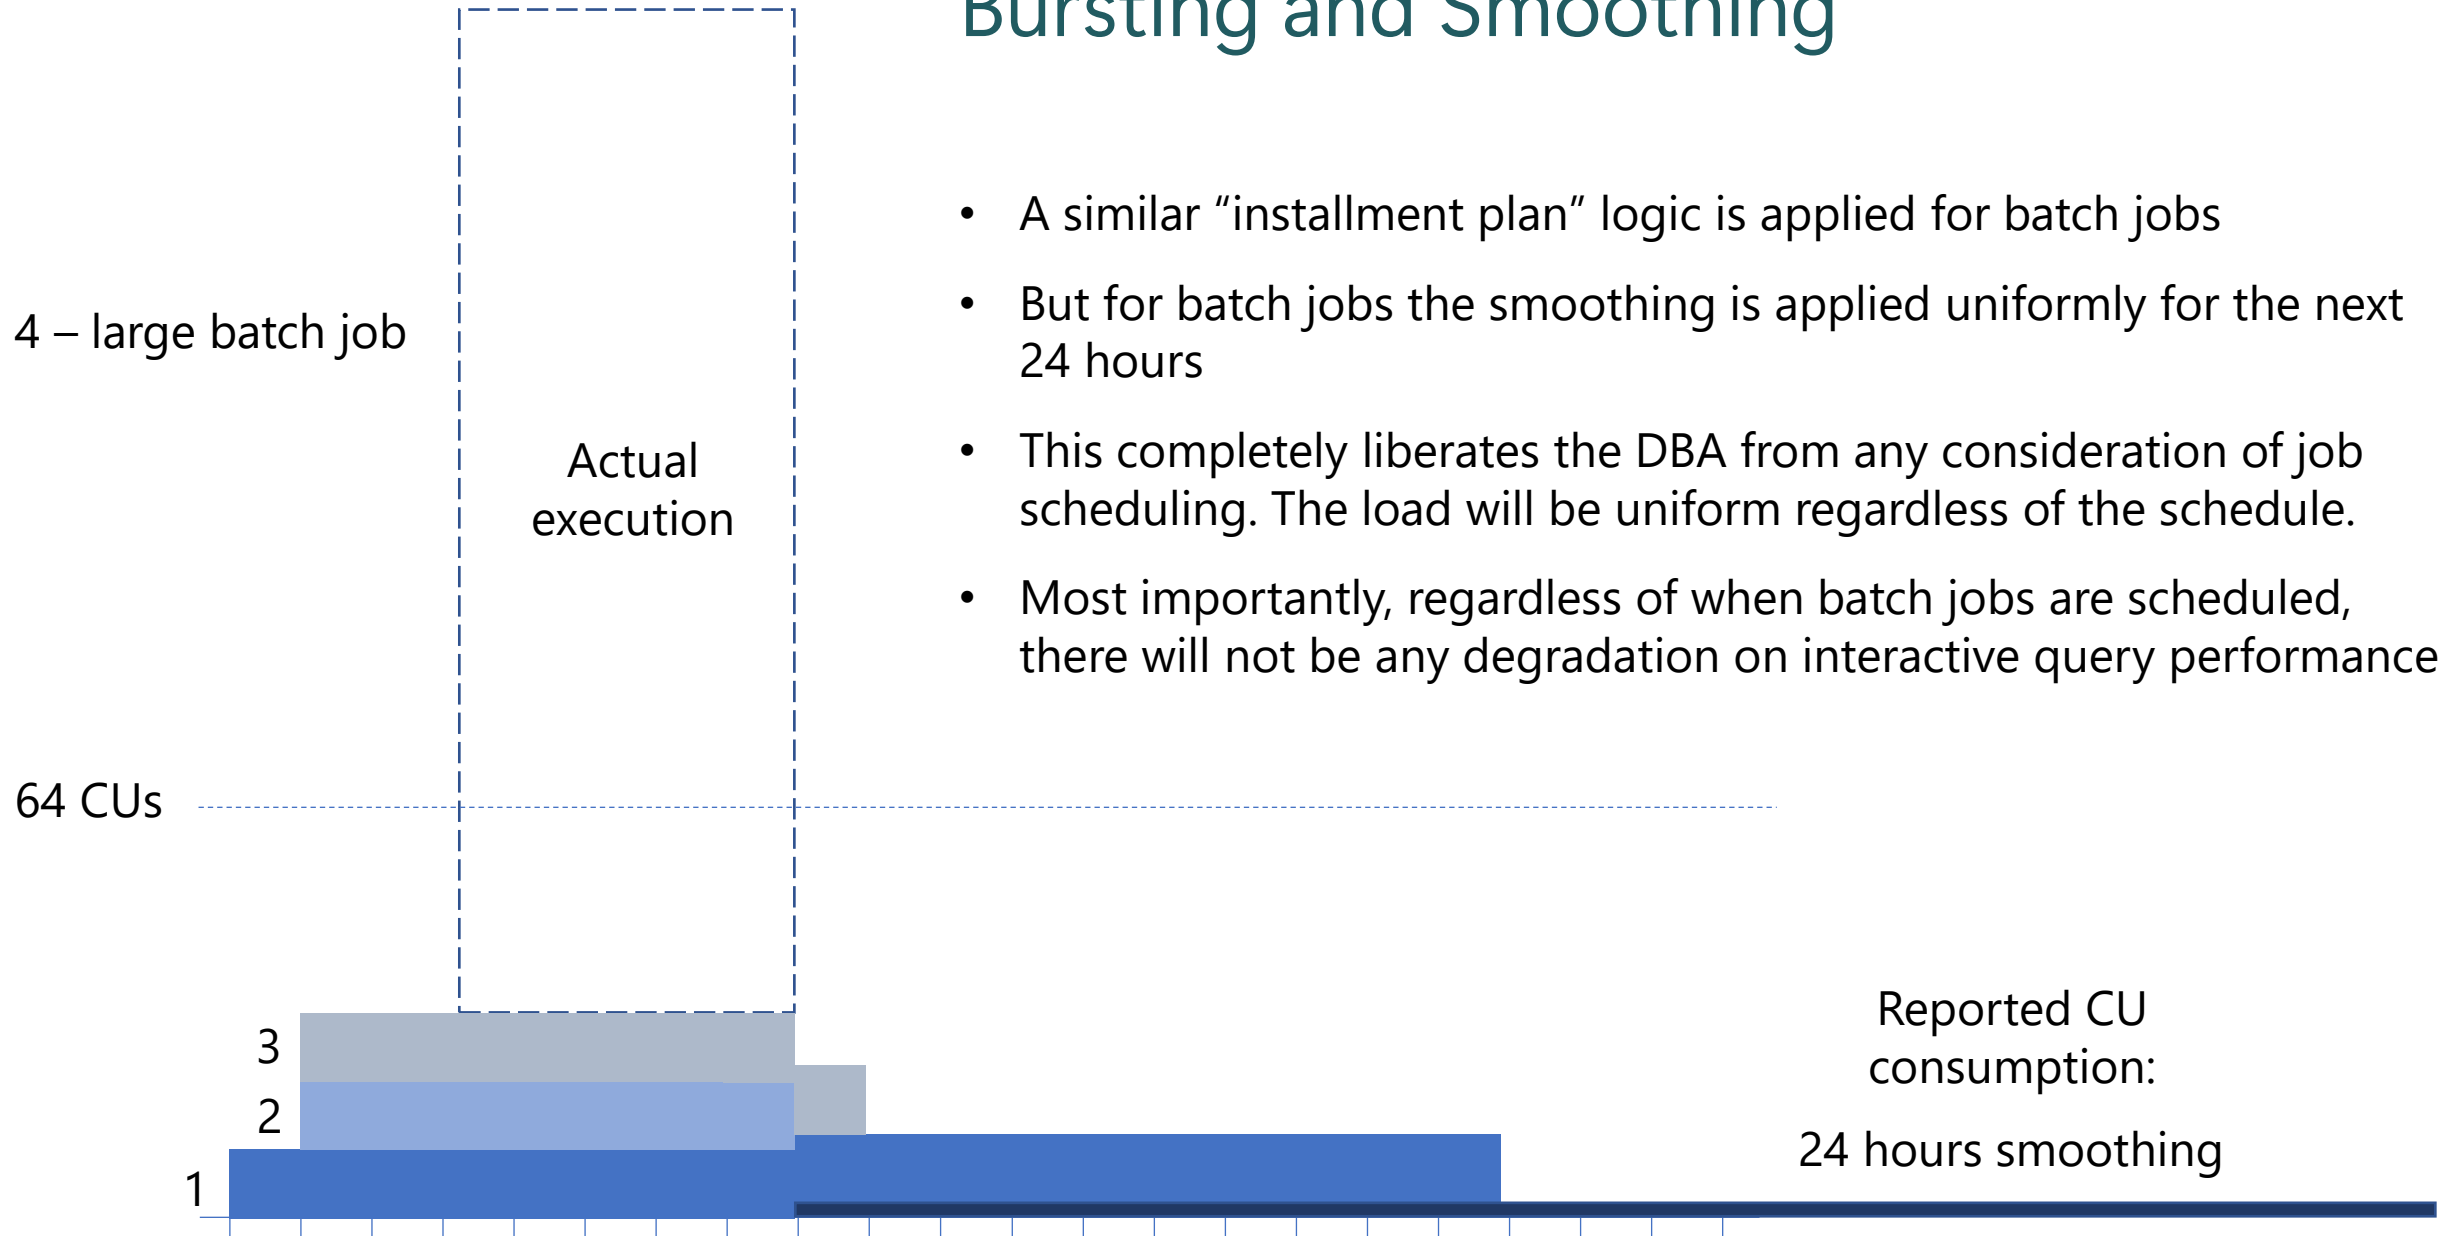
- No one (or few) queries can trigger an overload
- Instead of allowing runaway queries to create a local overload, Fabric smooths the queries reported usage to future time windows
- Kind of “Buy now, pay in the future” installment plan

Bursting and Smoothing



- Large batch processes traditionally were a threat to interactive queries as they would overload the compute resource
- DBAs traditionally had to carefully schedule these jobs to off-hours to avoid interference with interactive user experiences

Bursting and Smoothing



- A similar “installment plan” logic is applied for batch jobs
- But for batch jobs the smoothing is applied uniformly for the next 24 hours
- This completely liberates the DBA from any consideration of job scheduling. The load will be uniform regardless of the schedule.
- Most importantly, regardless of when batch jobs are scheduled, there will not be any degradation on interactive query performance

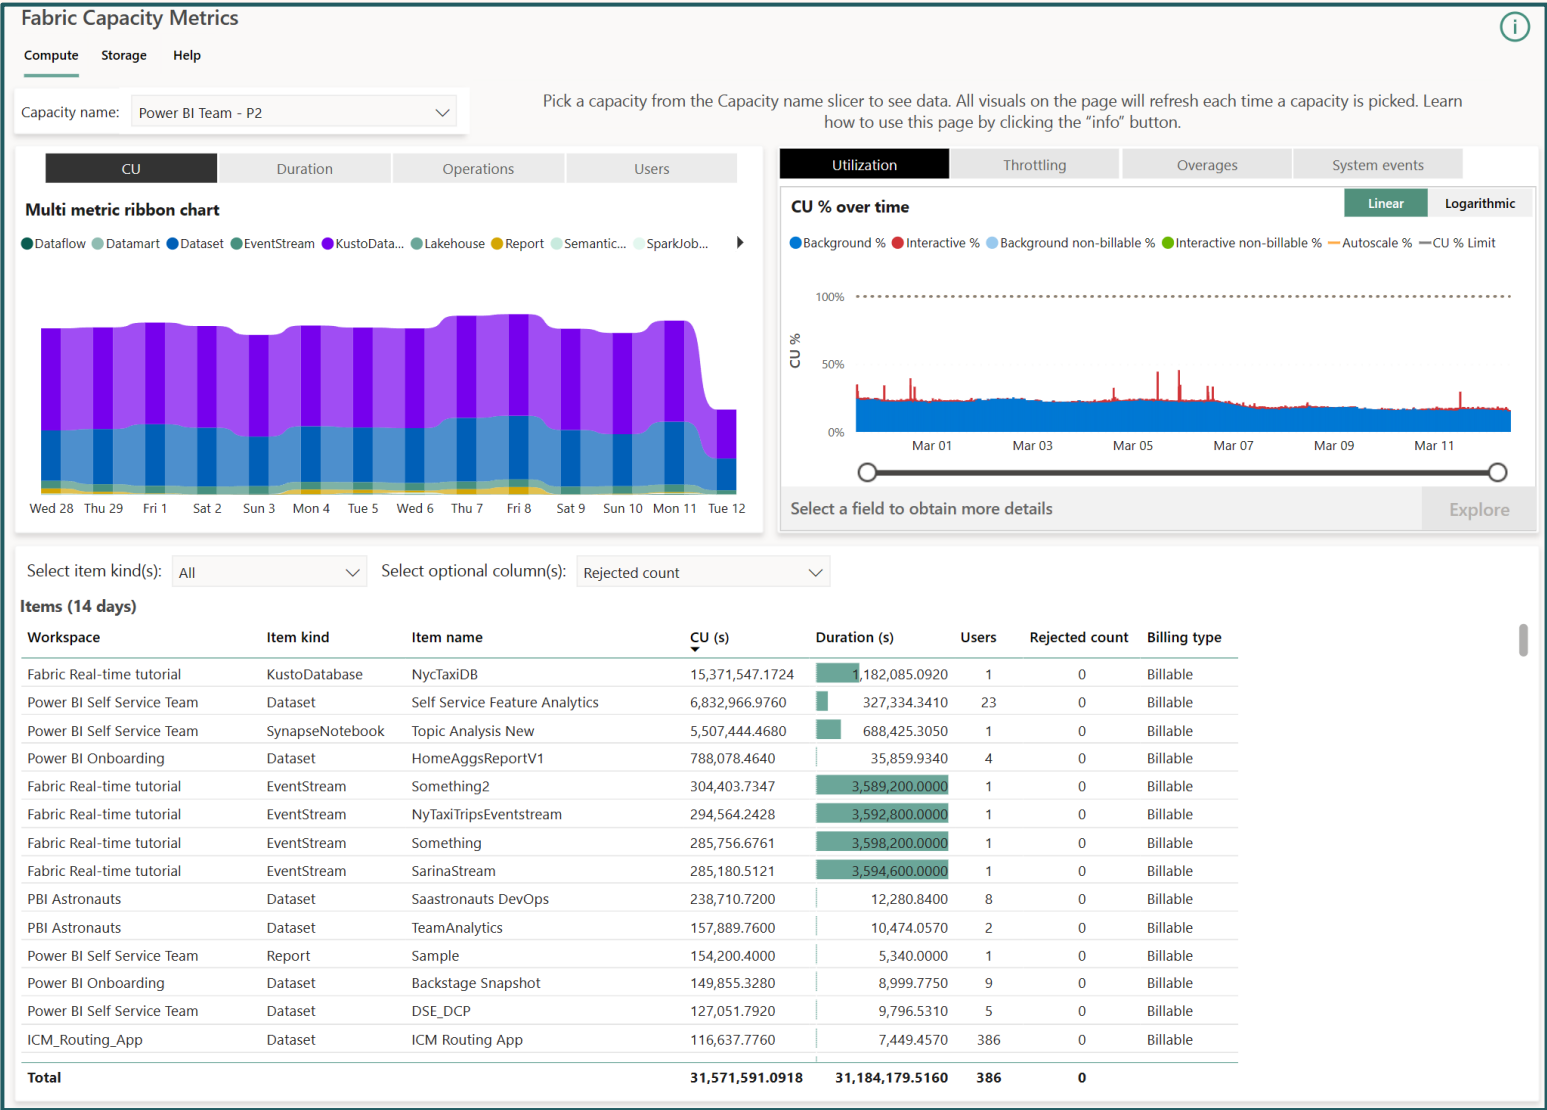


Monitoring with Capacity Metrics

Capacity Metrics

Monitor Capacities and Plan capacity scale-up with confidence

- Tenant wide visibility into capacity usage for all Fabric experiences
- Identify resource usage trends and their impact to autoscale & throttling
- View preview workload usage alongside production workloads to make data-driven capacity sizing decisions




Capacity Metrics

Monitor OneLake consumption

Measure the trends of workspace storage consumption against capacity limits, by day or hour

Reconcile costs with internal chargeback processes






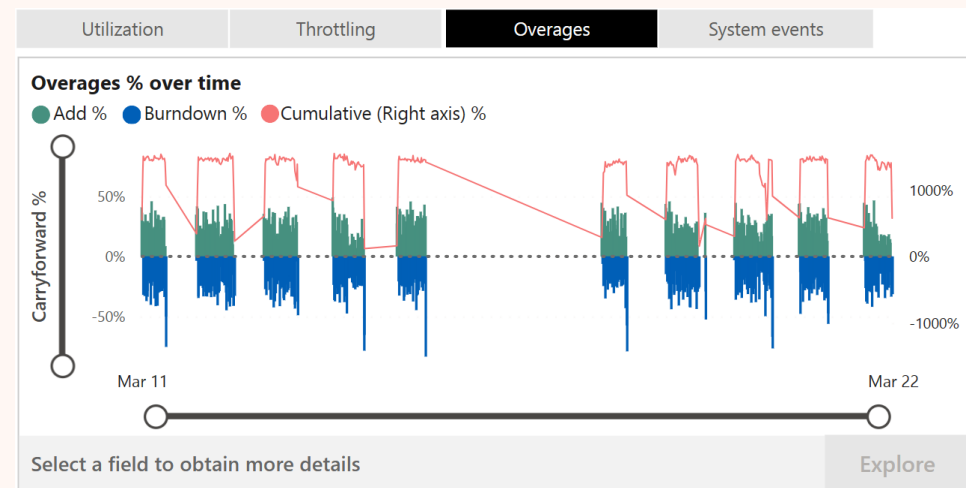



Capacity Throttling Policies

Throttling intro

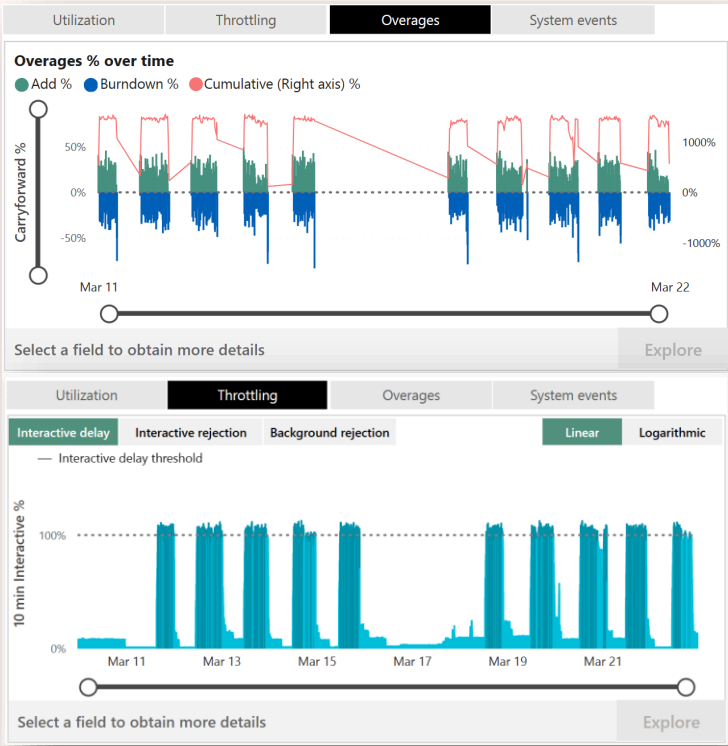
- Throttling is the platform policy for managing consumption that exceeds throughput is provided by SKU choice
- When workloads exceed the throughput of a capacity a cumulative debt is tracked to be burned down
- Cumulative debt is used to determine throttling policies and is burned down when resources are free

Overage Operation	Description
 Overages - Added	<ul style="list-style-type: none">• Timepoint when job requests exceed the throughput of a capacity, overages are added to the cumulative buffer to burn down.• This graph simplifies identification of the optimal timepoint to load timepoint drill to analyze the user operations that contributed to an overage.
 Overages - Burndown	<ul style="list-style-type: none">• Overages being reconciled when future capacity is free to burn down
 Overages - Cumulative	<ul style="list-style-type: none">• The total amount of queued work on the capacity to be burned down in the future when the capacity is not fully utilized



Capacity throttling evolution for Fabric

- For Fabric, throttling policies were refined to deliver multiple benefits
 - **Reduced throttling** for capacities that only experience occasional spikes
 - **Added overage protection** – rejection policies prevent overloaded capacities from irrecoverable overload
 - **Optimizations for long-running jobs:** We're optimizing the platform for long-running jobs, so if a job exceeds capacity limits, it will run to completion and the overage will be burned down against future capacity



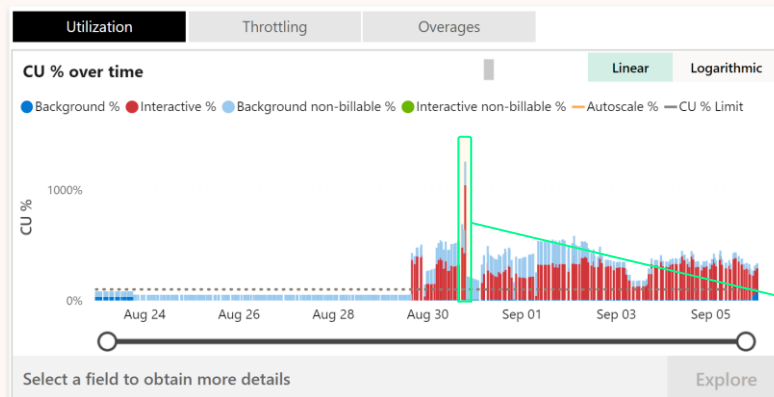
Smoothed Capacity - Future Use	Platform Policy	Customer Impact
$\leq 10m$	Overage Protection	Jobs can consume 10 minutes of future capacity use without throttling
$> 10m \rightarrow \leq 60m$	Interactive Delay	User requested interactive type jobs will be throttled
$> 60m \rightarrow \leq 24h$	Interactive Rejection	User requested interactive type jobs will be rejected
$> 24h$	Background Rejection	User Scheduled background jobs will be rejected from execution



Capacity Planning with Capacity Metrics

Capacity planning case study - measurement

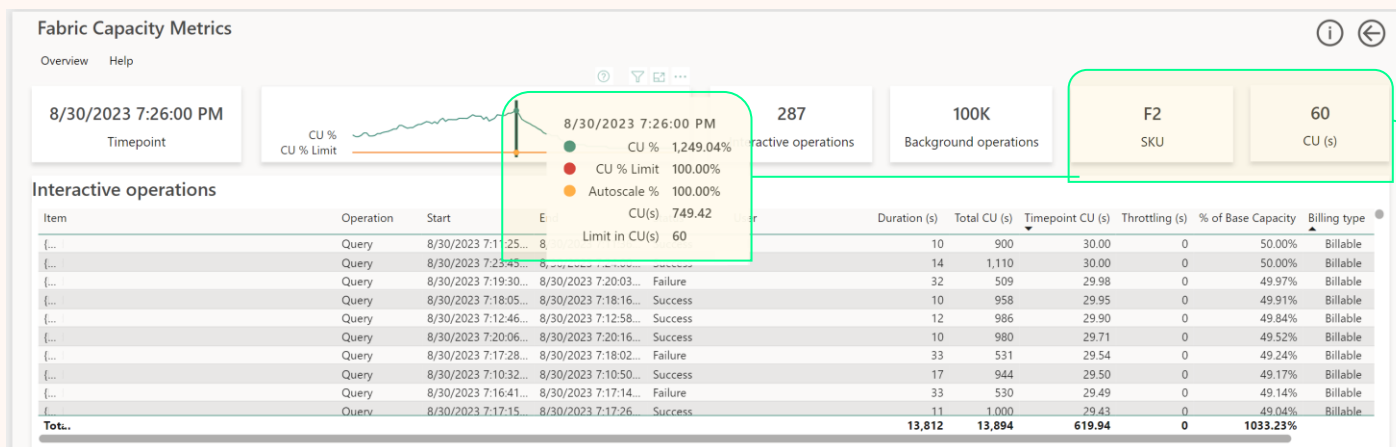
Start with a test or trial capacity to evaluate the load of specific Fabric Experiences i.e., Power BI Datasets, Spark Notebooks or a Datawarehouse



TimePoint	8/30/2023 7:26:00 PM
Background non-billable %	211.49%
Autoscale %	100.00%
Total CU Usage %	1250.53%
Total CU(s)	750.32
Interactive CU(s)	620.18
Background CU(s)	3.25
Background Preview CU(s)	126.90
100% in CU(s)	60

Right-click to drill through

If usage is above the current capacity limits , choose the desired utilization rate to accommodate via capacity scale up



Load Capacity Metrics timepoint drill to analyze :

- Total CU's consumed : 749 CU(s)
- Capacity Size : (F2)
- CU(s) available on your capacity : 60 CU(s)

Capacity planning case study – SKU selection


Universal Compute Capacities SKU Sizing

SKU	Capacity Units (CU)	CU's (per 30s)	Power BI SKU	Power BI V-cores
F2	2	60	-	0.25
F4	4	120	-	0.5
F8	8	240	A1	1
F16	16	480	A2	2
F32	32	960	A3	4
F64	64	1920	P1	8
F128	128	3840	P2	16
F256	256	7680	P3	32
F512	512	15360	P4	64
F1024	1024	30720	P5	128
F2048	2048	61440	-	256

To accommodate a **749 CU(s)** **load** the admin can purchase an F32 capacity providing 960 CU(s) of throughput



Pausing and Resuming Capacities



Introduction to Pausing and Resuming Capacities

Overview and Benefits

Pause and Resume lets you manage compute costs on F SKU capacities by suspending the execution of all workloads running on the capacity

- When a capacity administrator pauses a capacity:


Workloads stop
execution

New requests are
not run


Smoothed usage
will be reconciled

(details in the next demo)

Note: OneLake storage will remain active and billable while a capacity is paused




Bursting and Smoothing



- Job execution in Fabric workloads happens on-demand via capacity powered compute engines
- Fabric *bursting* will automatically allocate resources as needed to execute at maximum performance

Bursting and Smoothing

Actual
execution



- The Fabric capacity platform *smooths* usage out to reduce throttling which can occur when demand exceeds the throughput of the capacity that was purchased


Smoothing and Paused Capacities

Pause event on
Capacity



- When a capacity is paused...

64 CUs




Smoothing and Paused Capacities


Pause event on
Capacity



64 CUs



- When a capacity is paused...
- Usage that was smoothed into the future will be “reconciled” and charged against the capacity at the timestamp the capacity was paused
- Reconciled usage will show up as a spike in capacity metrics




Smoothing and Paused Capacities

Pause event on
Capacity



64 CUs



- When a capacity is paused...
- Usage that was smoothed into the future will be “reconciled” and charged against the capacity at the timestamp the capacity was paused
- Pause events can be viewed in the new System events tab

Utilization	Throttling	Overages	System events
System events			
State transition time	Capacity state	Capacity state change reason	
12/13/2023 9:12:14 AM	Active	Created	
12/13/2023 9:29:12 AM	Suspended	ManuallyPaused	
12/13/2023 9:30:15 AM	Active	ManuallyResumed	
12/13/2023 9:33:29 AM	Suspended	ManuallyPaused	
12/13/2023 9:34:58 AM	Active	ManuallyResumed	
12/13/2023 9:35:11 AM	Suspended	ManuallyPaused	
Select a field to obtain more details			Explore

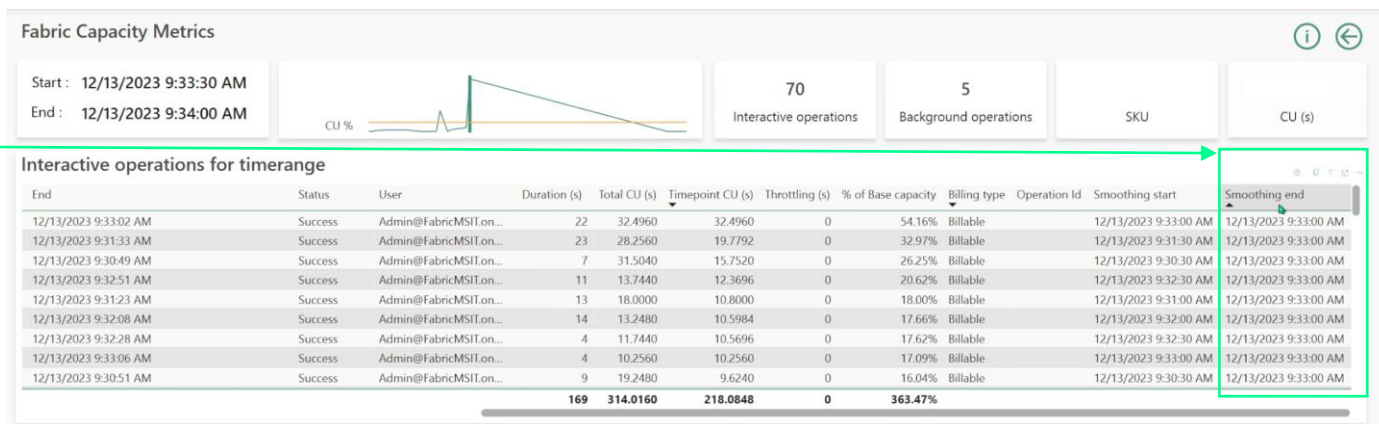
Smoothing and Paused Capacities

Pause event on
Capacity



64 CUs

- When a capacity is paused...
- Usage that was smoothed into the future will be “reconciled” and charged against the capacity at the timestamp the capacity was paused
- Pause events timestamp is shown in the smoothing end field in timepoint drill views





Bonus: Tips and Tricks
for capacity
management and
monitoring

"Fabric leverages the Power BI capacity model to bring even more powerful tools to your organization"



Citizen Devs



BI Specialists




Data Scientists



Data Engineers



DW/SQL
Professionals



Capacity Admins




Business
Leaders

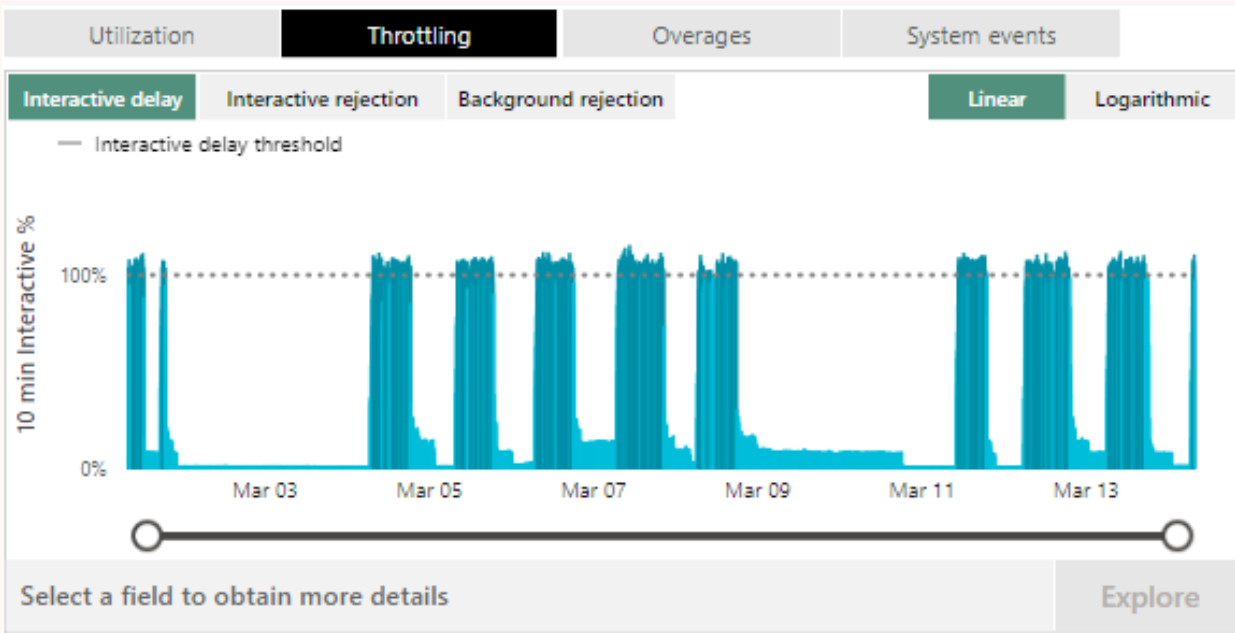


Employees
(Content Consumers)

My capacity is being throttled! What can I do?



Over 100% utilization doesn't always result in throttling

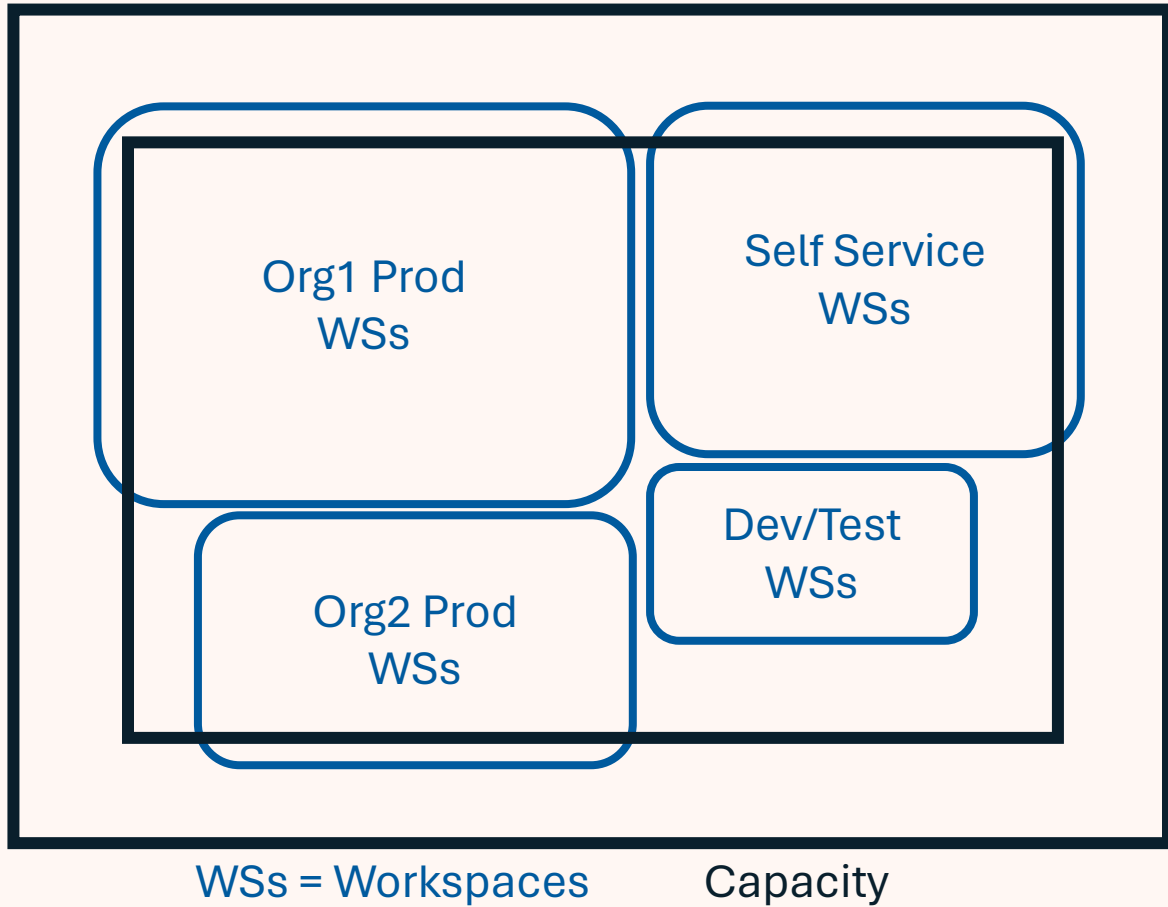


No penalty until you hit 100% on one of the throttling tabs

Note: For F SKU, if throttled, you can pause/resume to pay now and clear the carry forward, but that is not a long-term solution

When Capacity Units Run Out

Option 1 – Scale Up



Options to add compute

- Move to a bigger P SKU or RI F SKU
- Turn on autoscale (P SKU)
- Manual/Dynamic change size (F SKU)

Pros


- Add CUs for all items
- Easy

Cons

- Cost
- Bad actors (items with unintentionally high CU burn) can still be a problem

When Capacity Units Run Out

Option 2 – Scale Out



WSs = Workspaces

Capacity

Options

- Create multiple smaller P or F SKUs based on organization, type of work, etc.

Pros


- Easy
- Provides some isolation from bad actors (items with unintentionally high CU burn)
- Flexibility in capacity settings/governance

Cons

- Cost
- High CU items have increased chance of throttling

When Capacity Units Run Out

Option 3 – Optimize



WSs = Workspaces

Capacity

Approach

- Work with content creators to follow best practices and reduce CU consumption

Pros

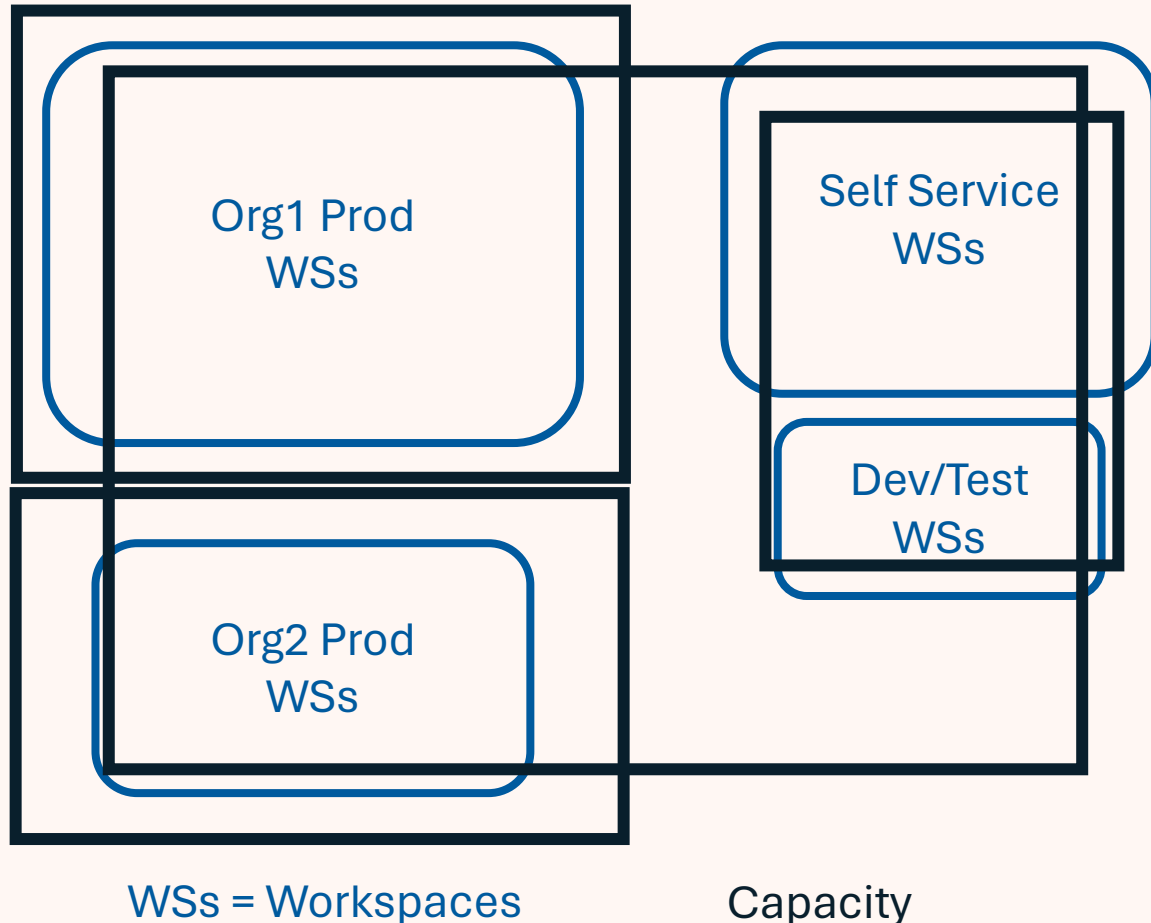
- Avoids increased cost
- Learning carries over to future content

Cons

- Can be difficult/time consuming

When Capacity Units Run Out

Option 4 – Isolate



Approach

- Provide isolated capacity for key items built by experienced developers

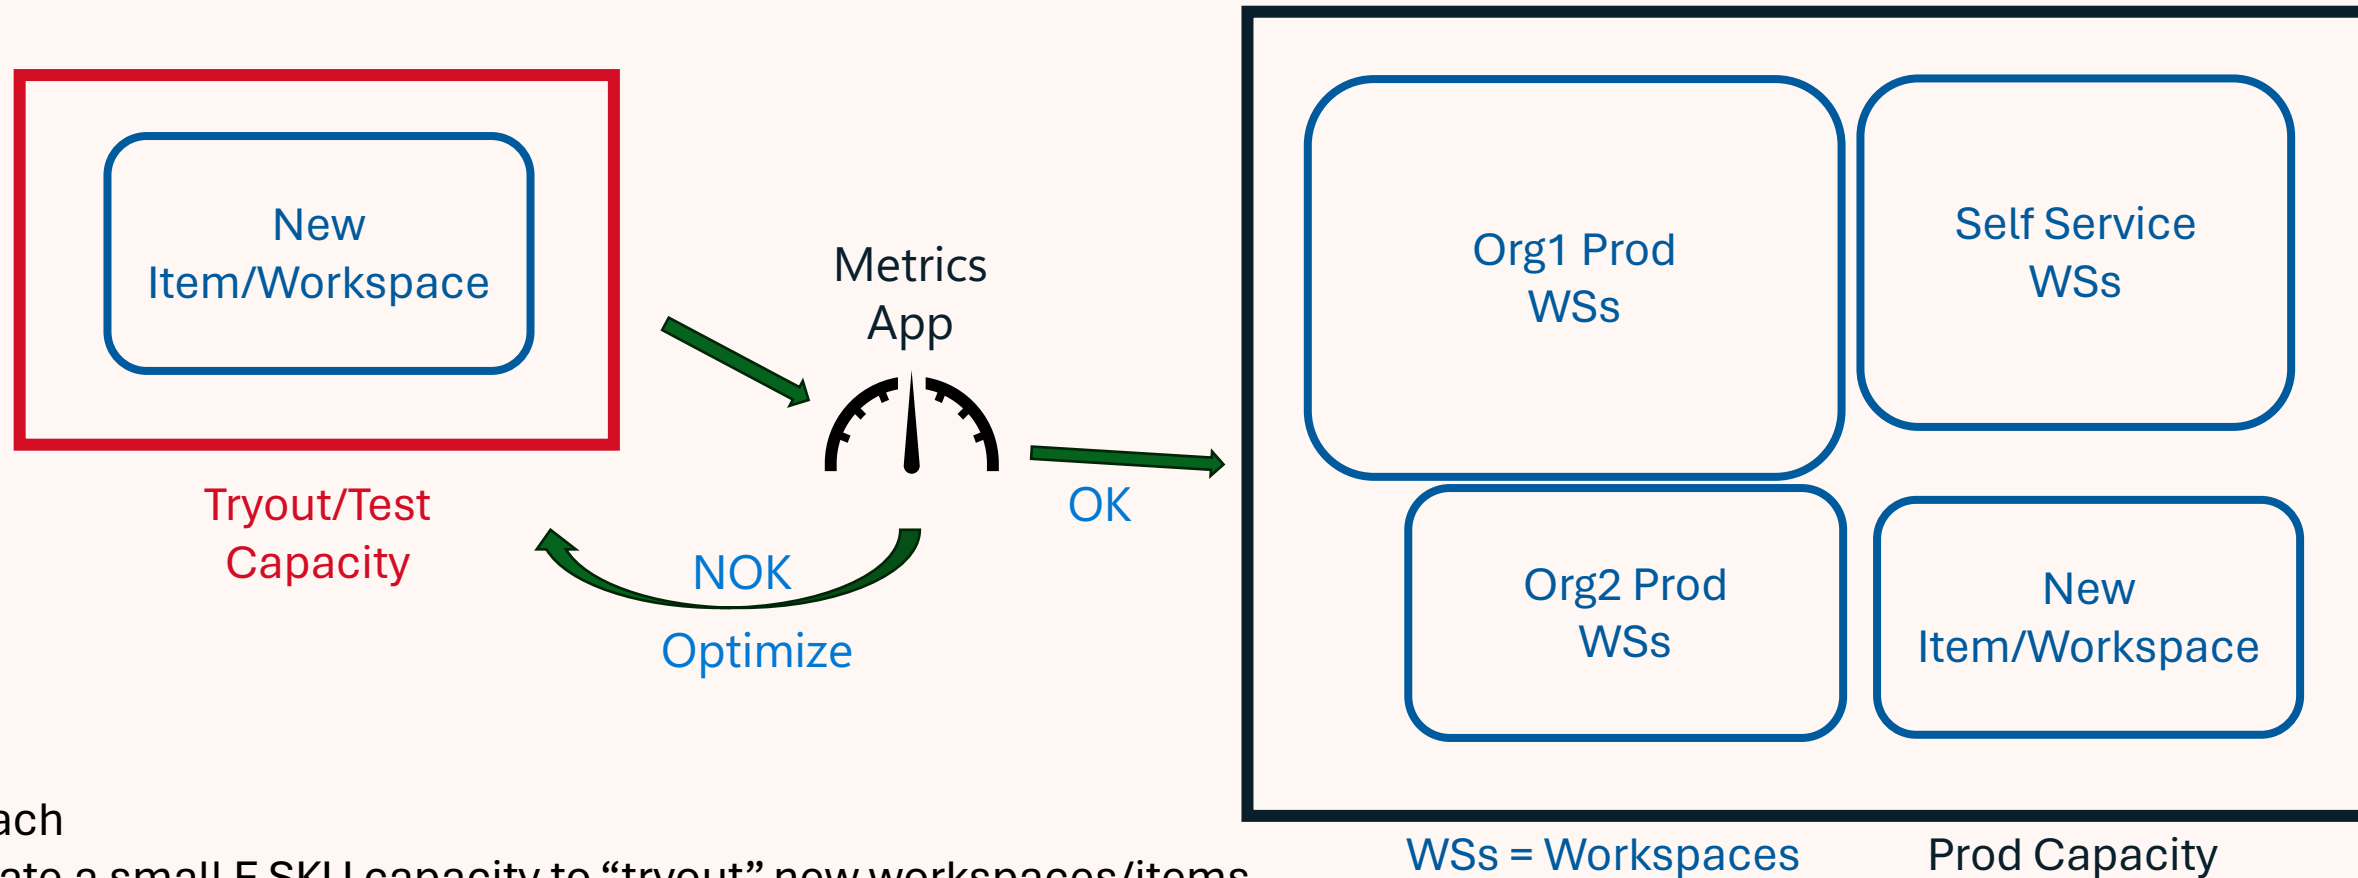
Pros

- Easy
- Provides isolation from items built by inexperienced developers and/or rapid unplanned usage growth
- Flexibility in capacity settings/governance

Cons

- Cost
- May lead to frustration of lower priority content developers/consumers

Isolation Strategy #4a – Tryout Capacity




Approach

- Create a small F SKU capacity to “tryout” new workspaces/items
- Assess CU consumption using metrics app
- If acceptable, move to prod capacity
- If not, optimize
- Pause tryout capacity when not in use, if possible
- Note size limits for semantic model size

Isolation Strategy #4b – Timeout Capacity

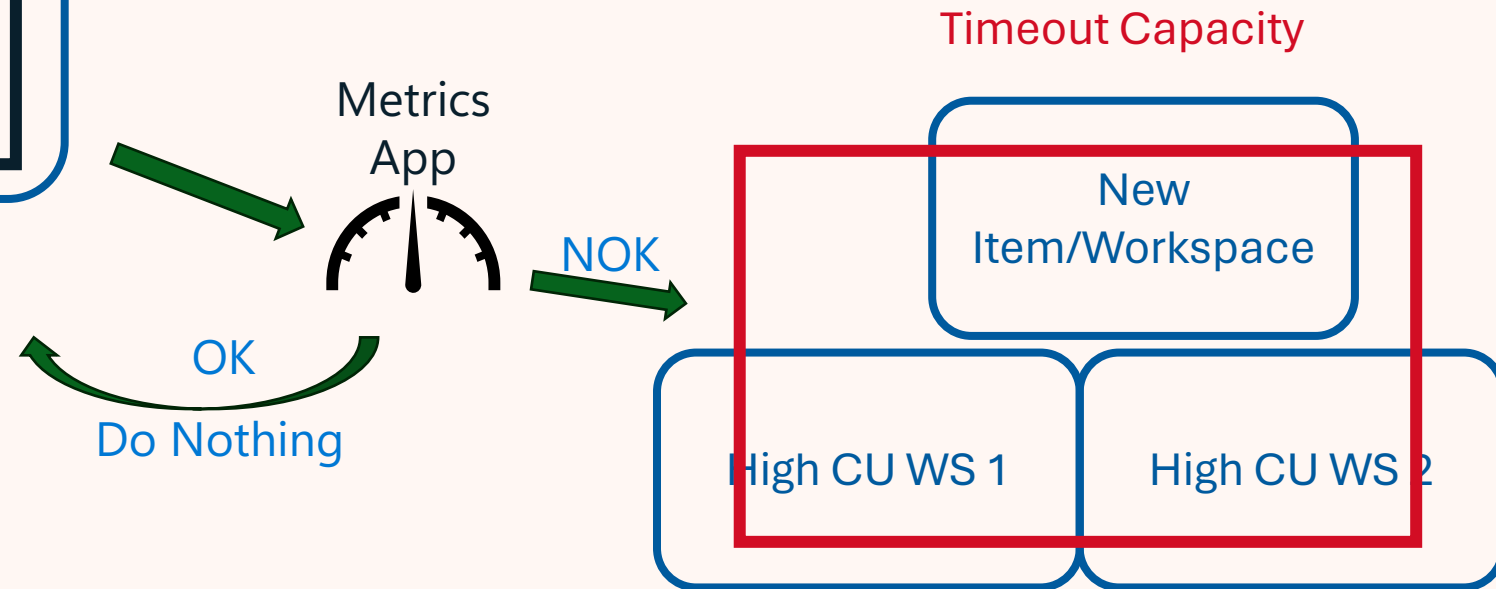
Approach

- Create a small F SKU capacity
- Assess CU consumption using metrics app
- If CU for new items/workspaces affects existing workloads (throttling), move WS to timeout capacity (Admin Portal/Capacity Settings)
- High CU items/WSs share smaller capacity (or you can pause it post move)
- Note size limits for semantic model size



WSs = Workspaces

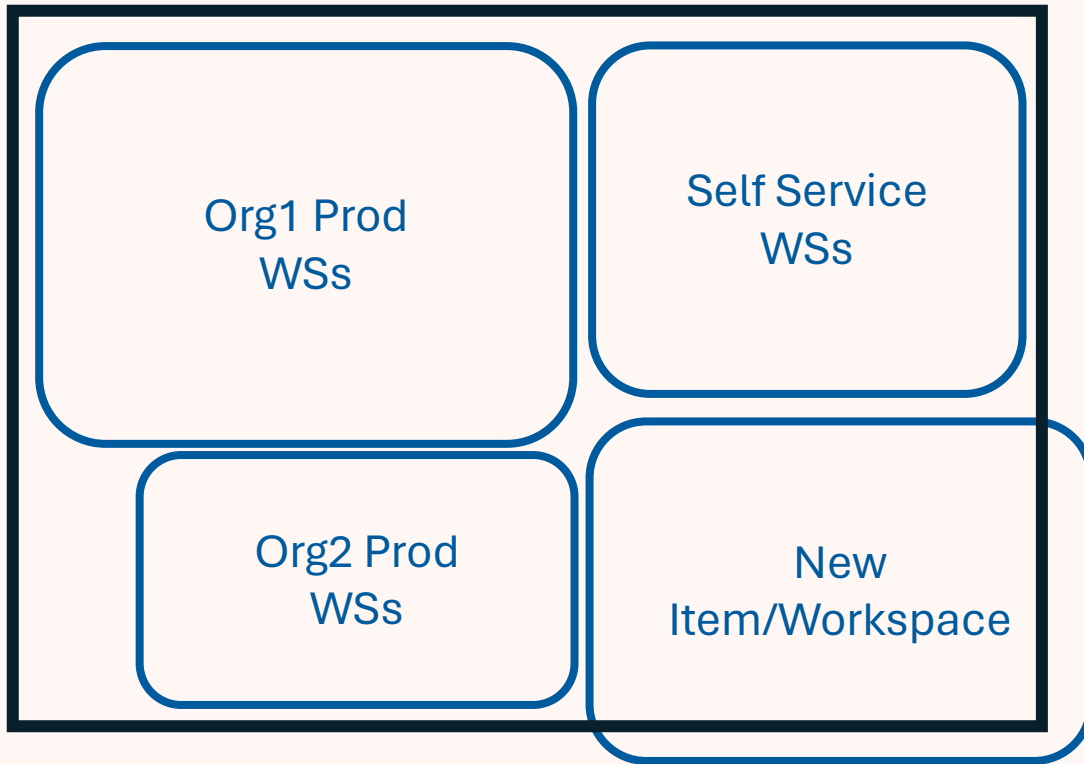
Prod Capacity



Isolation Strategy #4c – Rescue Capacity


Approach

- Create an F SKU capacity, keep it paused
- Assess CU consumption using metrics app
- If CU for new items/workspaces affects priority workloads (throttling), resume the new capacity and move priority WS to it (Admin Portal/Capacity Settings)
- Address issues with new content, then bring it back to original capacity, and pause the new one
- Note size limits for semantic model size



WSs = Workspaces

Prod Capacity



Rescue Capacity

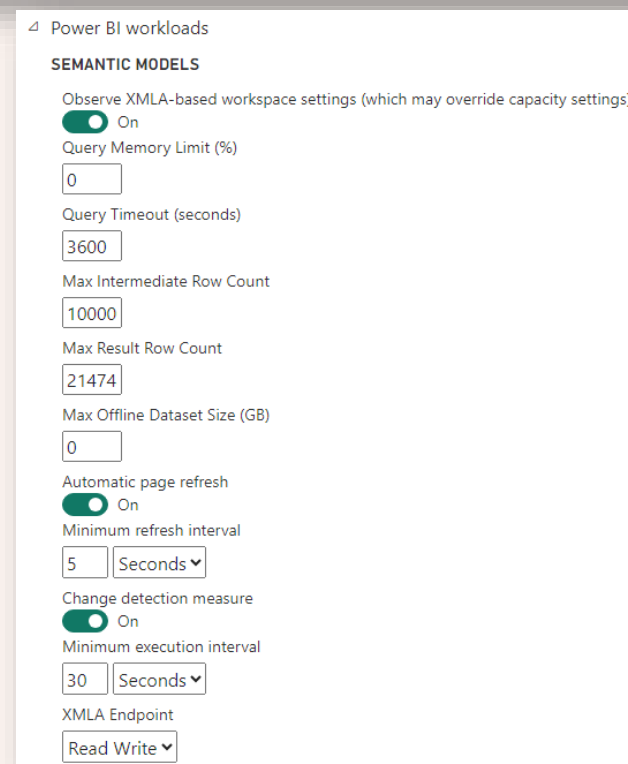
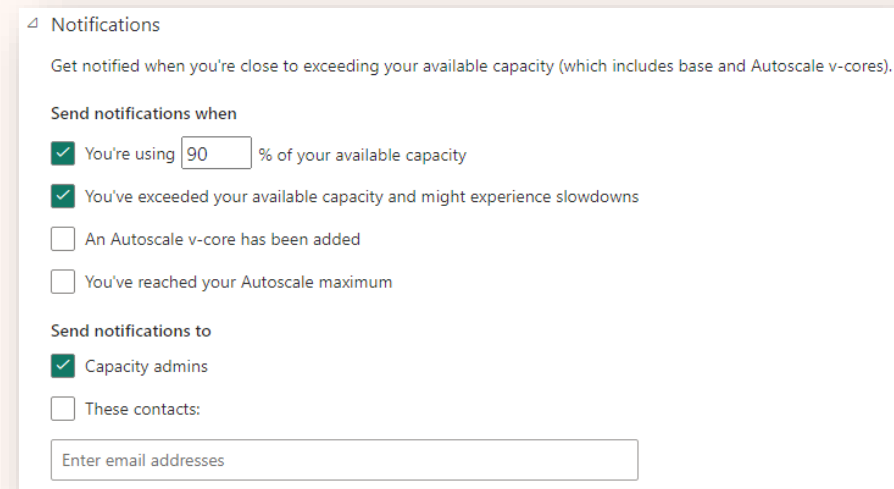



Recommendations for Cost/CU Savings

- Invest in education, knowledge/best practice sharing, COE, etc. for creators and consumers (proactive optimization)
- Avoid data/report sprawl (leverage certified/promoted models, OneLake shortcuts, etc.)
- Leverage a multi-capacity strategy (isolate, tryout, timeout, etc.)
- Right size your capacities and leverage F SKUs for pause/resume/resize, or reserved instances for discounts
 - Consider a combo of RI and PAYGO (for predictable surge activity)
- Choose the right tool for the job and stay up to date on Fabric feature releases
 - High concurrency mode for notebooks

Leverage the capacity settings in the UI

- Notifications on CU overuse
- Power BI workloads settings (e.g., query limits, page refresh)



Custom Solutions

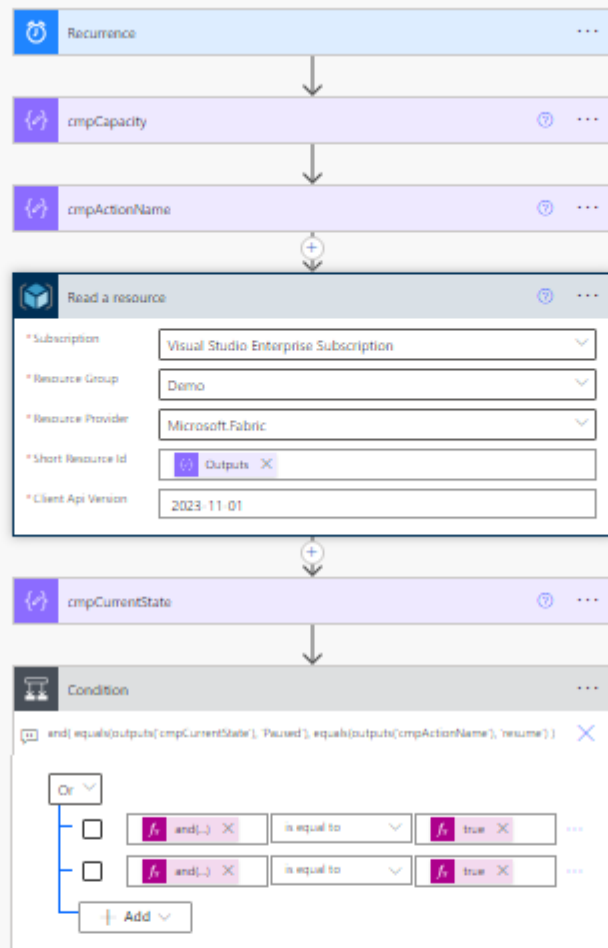
- Modify the Metrics App to meet your needs
- Build a custom report off the semantic model
- Send DAX queries to the metrics app semantic model in your own solution
 - Power Automate, Notebook (SemPy), PowerShell, etc.
 - Get throttling % values (Interactive Delay, Interactive Rejection, and/or Background Rejection)
 - Latest values and/or trends over time
 - Best for summarized data only (e.g., hour, day)

Collect data from multiple capacities and store it long term

Incorporate Metrics App queries into custom solutions

The screenshot shows the Metrics App interface. At the top, there is a purple bar labeled "DAX Query" with a question mark icon and a three-dot menu. Below this, there is a yellow bar labeled "Query" with a question mark icon and a three-dot menu. Under the "Query" bar, there are three input fields: "Workspace" with a dropdown menu showing "Custom value", "Dataset" with a dropdown menu showing "Custom value", and "Query text" with a text area containing "Outputs" and a close button. At the bottom left, there is a link "Show advanced options" with a dropdown arrow.

```
8 # Get max date from current delta table (to avoid loading duplicate days)
9 try:
10     df_max = spark.sql(f'''
11         SELECT MAX(Date) as MaxDate
12         FROM throttling;
13     ''')
14     maxdate = df_max.first()['MaxDate']
15 except:
16     maxdate = datetime.today() + timedelta(days=-6)
17     maxdateforDAX = maxdate.strftime('%Y,%m,%d')
18
19 if maxdate.date() < (datetime.today() + timedelta(days=-1)).date():
20
21     # Get data for each capacity, write daily csv and append delta
22     for capacity in lst_capacities:
23         querytext = '''\
24             DEFINE
25             MPARAMETER 'CapacityID' = "{capID}"
26             VAR yesterday =
27                 FILTER(ALL('Dates'[Date]), 'Dates'[Date] < TODAY() && 'Dates'[Date] > DATE({MD}))
28
29             EVALUATE
30             SUMMARIZECOLUMNS(
31                 'Dates'[Date],
32                 'TimePoints'[Start of Hour],
33                 yesterday,
34                 "IntDelay", ROUND( 'All Measures'[Dynamic InteractiveDelay %] * 100, 2 ),
35                 "IntReject", ROUND( 'All Measures'[Dynamic InteractiveRejection %] * 100, 2 ),
36                 "BackReject", ROUND( 'All Measures'[Dynamic BackgroundRejection %] * 100, 2 )
37             )
38         '''
39         df_throttling = fabric.evaluate_dax(workspace=MetricsWS, dataset=MetricsModel, dax_string=querytext)
40         if len(df_throttling) >= 1:
41             df_throttling.columns = df_throttling.columns.str.replace(r'(.*)[(\)](.*)', '', regex=True)
42             df_throttling.columns = df_throttling.columns.str.replace(' ', '_')
43             df_throttling['capacityId'] = capacity
44             filename = capacity + '_throttling_' + (datetime.today()).strftime('%Y%m%d') + '.csv'
45             df_throttling.to_csv("/lakehouse/default/Files/ThrottlingData/" + filename)
46             spk_throttle = spark.createDataFrame(df_throttling)
47             spk_throttle.write.mode("append").format("delta").option("overwriteSchema", "true").saveAsTable('Throttling')
```

Automate With F SKUs

- Pause/resume on a schedule
 - Automate with Power Automate, Logic Apps, or a Notebook
- Resize at peak/slow times
 - Mix with Reserved Instance (PAYGO when at increased size)
 - Query the metrics app and respond to actual demand (DIY autoscale)

Pause/Resume on
a Schedule

DIY Autoscale – Fabric Notebook

(Bret Myers)

Set SKU Ranges and Values

```
1 # Parameters to be passed in from pipeline.
2 minSku = 'F2' # min sku size we can scale down to
3 maxSku = 'F128' # max sku size we can scale up to
4 utilizationTolerance = 90 # Percentage of CU used to st
5 capacityName = 'fabricbamdemo' #capacity name to be mon
6 subscriptionId = '...'
7 metricsAppWorkspaceName = 'WS_FabricCapacityMetrics' #
8 metricsAppModelName = 'Fabric Capacity Metrics' # name
9 alertEmail = '' # email address to send alert that we s
```

Get credentials

```
1 tenantId = mssparkutils.credentials.getSecret('keyVaultEndpoint', 'secretName_tenantId')
2 clientId = mssparkutils.credentials.getSecret('keyVaultEndpoint', 'secretName_clientId')
3 secret = mssparkutils.credentials.getSecret('keyVaultEndpoint', 'secretName_clientSecret')
4
5 api_pbi = 'https://analysis.windows.net/powerbi/api/.default'
6 api_azuremgmt = 'https://management.core.windows.net/.default'
```

Not all code shown

FabricTools/CapacityAutoScale at main · bretamyers/FabricTools · GitHub

Query metrics app model

```
1 from azure.identity import ClientSecretCredential
2 import requests, json, math
3 from pyspark.sql.functions import explode
4
5 auth = ClientSecretCredential(tenant_id=tenantId, client_id=clientId, client_secret=secret)
6 access_token = auth.get_token(api_pbi)
7
8 header = {'Authorization': f'Bearer {access_token.token}', 'Content-type': 'application/json'}
9
10 body = {
11     "queries": [
12         {
13             "query": f"""
14                 DEFINE
15                     MPARAMETER 'CapacityID' = "{capacityId}"
16
17                     VAR __DS0FilterTable =
18                         FILTER(
19                             KEEPFILTERS(VALUES('TimePoints'[TimePoint])),
20                             'TimePoints'[TimePoint] >= NOW() - 1
21                         )
22
23                     VAR __DS0FilterTable2 = TREATAS({{"{capacityId}"}}, 'Capacities'[capacityId])
24
25                     VAR __DS0Core =
26                         SELECTCOLUMNS(
27                             KEEPFILTERS(
28                                 FILTER(
29                                     KEEPFILTERS(
30                                         SUMMARIZECOLUMNS(
31                                             'Capacities'[capacityId],
32                                             'Time'f'BillableTime'
```

Change SKU Size

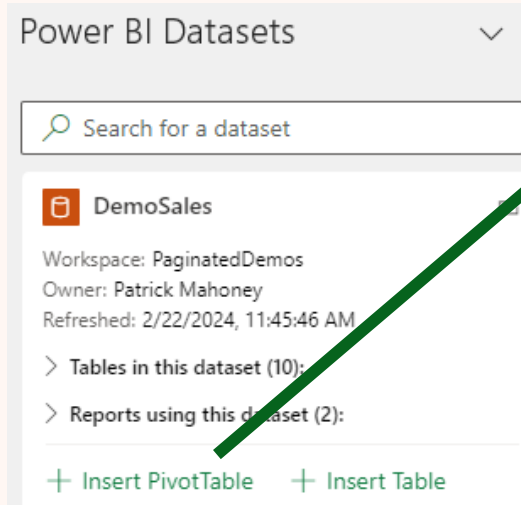
```
19 url = f'https://management.azure.com/subscriptions/{subscriptionId}/resourceGroups/{res
20
21 body = {
22     "sku": {
23         "name": f"{scaleSku}",
24         "tier": "Fabric"
25     }
26 }
27
28 response = requests.patch(url, headers=header, data=json.dumps(body))
```

Most Common Capacity Issues (Power BI)

Bad Practice	Recommendations/Typical Resolution
Model issues (M:M, bi-di, snowflake, etc.) and/or inefficient DAX	Follow best practices (e.g., BPA), star schema
Too many visuals	Multi card, small multiples, Deneb, PowerPoint background, etc.
Big single visual (i.e., matrix with lots of rows, columns, and/or measures)	Improve report design (e.g., drillthrough, apply all Slicers, report page tooltip), field parameters, calc group guardrails, etc.
Complex RLS	Remodel to enable simple filter like Table[Email] = USERPRINCIPALNAME()
Very high concurrency	Optimize reports, DAX, etc. (big multiplier) Consider QSO, data subsets
Direct Query	Switch to import or Direct Lake, if possible. Aggregations, hybrid tables, etc.
Analyze in Excel	Automate downstream analytics with a Power BI report instead, subscriptions, DAX connected table, slicers/measures first, etc.
Excessive refresh	Don't "break the fold", incremental refresh, reduce frequency, optimize M code

Save Those CUs – Getting Data Into Excel

Analyze in Excel



Connected Table

Key Takeaways

- How you build it matters
 - Filters & measures first!
- This shows durations but it's CU that matters (test your use cases/models)
- Opt for DAX Connected Tables
 - Create pivot table from that, if needed

✗ Rows, Measures, Filter

StartTime	Type	Duration	User	Database	Query
11:49:30	MDX	2,328ms	Power BI...	DemoSales	SELECT {[Measures].[To
11:49:26	MDX	0ms	Power BI...	DemoSales	SELECT {AddCalculated
11:49:23	MDX	0ms	Power BI...	DemoSales	SELECT {AddCalculated
11:49:17	MDX	1,875ms	Power BI...	DemoSales	SELECT {[Measures].[To
11:49:03	MDX	4,469ms	Power BI...	DemoSales	SELECT {[Measures].[To
11:48:54	MDX	3,938ms	Power BI...	DemoSales	SELECT {[Measures].[To

✓ Filter, measures, rows

StartTime	Type	Duration	User	Database	Query
10:06:13	MDX	1,625ms	Power BI...	DemoSales	SELECT {[Measu
10:06:03	MDX	781ms	Power BI...	DemoSales	SELECT {[Measu
10:05:49	MDX	109ms	Power BI...	DemoSales	SELECT {[Measu
10:05:46	MDX	312ms	Power BI...	DemoSales	SELECT {[Measu
10:05:43	MDX	234ms	Power BI...	DemoSales	SELECT FROM [M
10:05:14	MDX	0ms	Power BI...	DemoSales	SELECT {AddCal

Refresh (same for both)

StartTime	Type	Duration	User	Database	Query
11:50:30	MDX	2,234ms	Power BI...	DemoSales	SELECT {[Measures].[To

✗ Rows, Measure, Filter

StartTime	Type	Duration	User	Database	Query
01:28:50	DAX	31ms	Power BI...	DemoSales	DEFINE VAR __C
01:28:41	DAX	1,516ms	Power BI...	DemoSales	DEFINE VAR __C
01:28:40	DAX	16ms	Power BI...	DemoSales	DEFINE VAR __C
01:28:34	DAX	156ms	Power BI...	DemoSales	DEFINE VAR __C
01:28:33	DAX	16ms	Power BI...	DemoSales	DEFINE VAR __C
01:28:31	DAX	0ms	Power BI...	DemoSales	DEFINE VAR __C
01:28:30	DAX	141ms	Power BI...	DemoSales	DEFINE VAR __C
01:28:15	DAX	2,047ms	Power BI...	DemoSales	DEFINE VAR __C
01:28:11	DAX	1,797ms	Power BI...	DemoSales	DEFINE VAR __C
01:28:08	DAX	594ms	Power BI...	DemoSales	DEFINE VAR __C
01:27:56	DAX	281ms	Power BI...	DemoSales	DEFINE VAR __C
01:27:50	DAX	16ms	Power BI...	DemoSales	DEFINE VAR __C

✓ Filter, measures, rows

StartTime	Type	Duration	User	Database	Query
09:14:20	DAX	16ms	Power BI...	DemoSales	DEFINE VAR __DS0Filt
09:14:07	DAX	1,000ms	Power BI...	DemoSales	DEFINE VAR __DS0Filt
09:14:02	DAX	1,188ms	Power BI...	DemoSales	DEFINE VAR __DS0Filt
09:13:59	DAX	594ms	Power BI...	DemoSales	DEFINE VAR __DS0Filt
09:13:51	DAX	531ms	Power BI...	DemoSales	DEFINE VAR __DS0Filt
09:13:50	DAX	0ms	Power BI...	DemoSales	DEFINE VAR __DS0Cor

Refresh (same for both)

StartTime	Type	Duration	User	Database	Query
11:54:49	DAX	1,969ms	Power BI...	DemoSales	DEFINE VAR __DS0FilterTable = TREATA

I PITY THE FOOL

**THAT DOESN'T
GIVE SESSION FEEDBACK**



Slides



https://github.com/BenniDeJagere/Presentations/{Year}/{YYYYMMDD}_{Event}





Thank you

